

Political Conversations with Claude

How Much, Where, and Whether Elections Move Them

Andrew B. Hall

Daniel M. Thompson

July 28, 2026

Abstract

Using the Anthropic Economic Index’s public state-month topic shares, we document how much of Claude usage is political, what people are doing when they raise politics, and whether elections move it. Politics is a small share of all conversations—0.46% in May 2026, or roughly one conversation in 217—and is overwhelmingly personal rather than work-related. It is also overwhelmingly *informational*: 67.6% of political conversations produce an explanation or an analysis, while only 2.8% produce a message, a plan, or organizing content—a ratio of about 24 to 1. People are using Claude to understand politics, not yet to act in it. We then ask whether holding a primary election moves the political share. States with May 2026 primaries saw a 0.10 percentage-point increase off a 0.48% base, about 22%. The effect is specific: it appears for politics and not for news, geopolitics, or a non-political comparison topic. Applying the identical specification to all 210 topics in the release, politics ranks 3rd largest, giving a randomization-inference p -value of 0.024.

1. Data

All figures come from the Anthropic Economic Index (AEI) release of 26 June 2026 (CC-BY, [Hugging Face](#)), which reports the share of Claude conversations assigned to each request topic for April and May 2026, at global, country, and US state level. Two features of the data govern everything that follows. First, shares are rounded to two decimals, so the measurement grid is coarse relative to the effects at issue. Second, a missing state-topic cell means the cell fell below Anthropic’s privacy threshold, not that the share was zero. Topic taxonomies are re-clustered in every release, so these shares are not comparable to earlier or later AEI releases.

Primary dates are hand-compiled from NCSL and FVAP, which agree; the file and its sources are documented in `original_data/primary_dates/`.

1.1. How Anthropic codes conversations

Because everything below depends on it, it is worth being precise about how these categories are produced. All AEI classification is done by [Clio](#), Anthropic’s privacy-preserving analysis system: conversation transcripts are read only by another instance of Claude, never by a human, and

cells with too few observations are suppressed before release. Every variable we use is therefore a model-assigned label aggregated over sampled conversations, not a human-coded one.

The topic hierarchy. Requests are grouped into a three-level topic tree. Level 0 (“Detailed”) has 1012 topics, level 1 (“Minor”) has 196, and level 2 (“Major”) has 20. Our outcome, “Politics and public record,” is a *Minor* topic, while “Politics” and “Elections” are Detailed topics. The published file gives each node’s level but not its parent, so while the levels are ordered by aggregation we cannot confirm from the data which Detailed topics roll up into which Minor one. Critically, this taxonomy is re-clustered from scratch in every release, so a topic’s share cannot be compared across releases even under the same name.

Use case assigns each conversation to **work**, **personal**, or **coursework**.

Collaboration pattern assigns each conversation to one of five modes (defined in Panel C of Table 2). These are then collapsed into two buckets: *automation* (directive or feedback loop) and *augmentation* (task iteration, learning, or validation). One subtlety matters for reading the numbers: the buckets renormalize over conversations that received a pattern, excluding the **none** category, so bucket shares sum to 100 while pattern shares sum to 100 only once **none** is included. We verified this against the data—directive plus feedback loop, divided by one minus the **none** share, reproduces the published automation figure exactly—and the check is asserted in the code that generates this memo.

Artifact records the single most prominent concrete output the conversation produced, across 30-odd labels plus **none**. This is the variable we lean on hardest below, because it distinguishes a conversation that ended in an explanation from one that ended in a draft email.

A note on reading the tables. Because these classifications are nested in some places and merely parallel in others, every table below that mixes them is split into labelled panels. Rows within a panel are mutually comparable and, where noted, sum to 100. Rows in different panels are *not* comparable and must never be added together. In particular, AEI does not publish the parent-child map linking Detailed topics to their Minor parents, so we never assert that a given Detailed topic sits beneath a given Minor one.

2. Descriptives

2.1. Politics is a small share of conversations

In May 2026, 0.46% of US Claude conversations were assigned to “Politics and public record,” up from 0.42% in April. That is roughly one conversation in 217. The narrower “Elections” topic is smaller still, at 0.07% in May. Any claim about AI and political information should be read against this base: political conversation is a rounding error in aggregate Claude usage, before and after any election.

Table 1 decomposes the political topics, split by taxonomy level. Panel A holds the Minor topics—aggregates, one of which is our outcome—and Panel B the Detailed leaves. Keeping

them apart matters: “News aggregation” (0.49% in May) is numerically larger than “Politics and public record” (0.46%), but it is a Detailed topic and the latter is a Minor one, so the comparison is between units at different levels of aggregation rather than between a category and its sub-category. The rightmost columns show how few states clear the privacy threshold for the narrower topics, which is why the analysis below uses the broad politics measure as its outcome.

Table 1: Political and News Topics: National Shares and State Coverage

	US National Share (%)		# States Reported	
	April	May	April	May
<i>Panel A. Minor topics (level 1, aggregates)</i>				
Politics and public record	0.42	0.46	36	35
Geopolitics and strategy	0.20	0.15	21	22
Government filings	0.06	0.06	7	7
<i>Panel B. Detailed topics (level 0, leaves)</i>				
News aggregation	0.42	0.49	29	34
News writing	0.15	0.16	20	24
Politics	0.14	0.15	19	25
Geopolitics	0.08	0.07	7	11
Elections	0.04	0.07	3	8
Public records lookup	0.05	0.05	4	12

2.2. The United States is the most political market

Across the 59 countries for which the politics share is reported in May 2026, the United States ranks 1st at 0.46%, against 0.09% at the bottom (VNM)—a spread of roughly $5.1\times$ (Figure 2). Political use of Claude is not uniform across markets, and the US sits at the top of the distribution rather than in the middle of it.

2.3. Political use is personal, not professional

Table 2 reports how political conversations are conducted. The pattern is stark: 78.4% of politics conversations are personal use, against only 18.9% work use and 2.7% coursework. Whatever people are doing when they ask Claude about politics, they are mostly not doing their jobs. Interaction style is majority augmentation (58.3%) rather than automation (41.7%), and both shift slightly toward augmentation between April and May.

A caveat that constrains this section: these behavioral metrics are published *only* at the global level. Country and state cells carry the topic share alone. The composition results are therefore descriptive and cannot be linked to the state-level variation used below.




Figure 1: Political topic shares, April versus May 2026. US national shares of all Claude conversations. Domestic political topics rise between April and May while geopolitical topics fall.

Table 2: How Political Conversations Are Conducted (Global, %)

Definition		April	May	Change
<i>Panel A. Use case (sums to 100)</i>				
Personal		77.7	78.4	0.7
Work		19.7	18.9	-0.8
Coursework		2.6	2.7	0.1
<i>Panel B. Collaboration bucket (sums to 100, excludes none)</i>				
Augmentation	Task iteration, learning, or validation	56.9	58.3	1.4
Automation	Directive or feedback loop	43.1	41.7	-1.4
<i>Panel C. Collaboration pattern (sums to 100; Panel B is built from these)</i>				
Learning	Seeking understanding	40.5	40.7	0.2
Directive	Minimal human interaction	35.0	33.9	-1.1
Task iteration	Human refines AI outputs	12.0	13.3	1.2
Feedback loop	Iterative dialogue with feedback	6.9	6.7	-0.2
Validation	Human checking own work	2.9	2.9	-0.0
None	No pattern assigned	2.7	2.6	-0.1

2.4. Seeking information about politics versus doing politics

The composition numbers hint at something the artifact data make unmistakable. Table 3 and Figure 3 compare what political conversations produce against the baseline across all Claude




Figure 2: Politics-topic share by country, May 2026. Twelve highest and twelve lowest of 59 reported countries.

conversations.

Political conversations are dominated by outputs that transmit understanding. 47.7% end in an explanation or answer, against 16.7% of conversations generally—nearly three times the rate. Another 19.9% end in an analysis or summary. Together, 67.6% of political conversations produce an explanation or an analysis.

Outputs that constitute political *action* are not merely rare in absolute terms; they are rarer than in ordinary Claude usage. Only 0.9% of political conversations produce an email or message—the artifact one would expect from contacting a representative, writing public comment, or organizing—against 4.6% of conversations generally. Plans or strategies account for 1.8%. Summing the action-shaped artifacts gives 2.8%, against 67.6% for the information-shaped ones: a ratio of roughly 24 to 1. The collaboration patterns tell the same story from a different angle, with 40.7% of political conversations classified as *learning*—seeking understanding—the single largest pattern (Table 2, Panel C).

2.5. Reading this through the three layers

Hall’s framework for AI and democracy distinguishes three layers at which the technology could matter: an *information* layer, in which AI makes citizens and governments smarter; a *representation* layer, in which AI acts on citizens’ behalf so that their interests are conveyed faithfully; and

Table 3: Artifacts Most Over- and Under-Represented in Political Conversations (Global, May 2026)

Artifact produced	Politics (%)	All conversations (%)	Difference
Explanation or answer	47.7	16.7	31.1
Analysis or summary	19.9	5.5	14.3
None	15.0	6.5	8.5
Script or snippet	0.2	3.0	-2.8
Educational material	0.0	2.8	-2.8
Code fix or debug	0.0	3.3	-3.3
Email or message	0.9	4.6	-3.7
App or website	0.3	4.2	-3.9
Advice or recommendation	3.3	10.7	-7.4
Document or report	6.6	14.9	-8.3




Figure 3: What political conversations produce. Share of conversations whose most prominent output is each artifact type, politics topic versus all conversations.

a *governance* layer, in which we set the rules under which the AI systems themselves operate.¹ The artifact distribution offers the first direct measurement we are aware of how present-day usage divides across them.

The headline is that the information layer dominates—but the aggregate hides a distinction

¹Andrew B. Hall, “Building Political Superintelligence,” *Free Systems*, March 26, 2026.

that matters for the framework, and Table 4 draws it out by running the same artifact split topic by topic.

Table 4: Information, Action, and Document Artifacts by Topic (Global, May 2026)

Topic	Level	Share (%)	Information	Action	Document
Politics	Detailed	0.15	74.4	1.1	3.3
Politics and public record	Minor	0.46	67.6	2.8	8.1
Public records lookup	Detailed	0.05	64.9	0.2	4.0
Elections	Detailed	0.07	62.4	1.2	6.2
Government filings	Minor	0.06	28.2	11.1	39.6

The civic-informational topics are overwhelmingly informational. “Politics” is 74.4% explanation-or-analysis, “Elections” 62.4%, and—despite its transactional-sounding name—“Public records lookup” 64.9%. Action artifacts are near zero across all three (1.2% for Elections, 0.2% for public records). A public-records lookup turns out to be exactly what it says: a lookup. This is Condorcet’s printing-press argument running on new hardware, and the primary-election result in Section 3 shows it responding to the electoral calendar exactly as a demand-for-information story predicts.

“Government filings” is the exception, and it is a different kind of thing entirely. Only 28.2% of those conversations produce an explanation or analysis. Instead 39.6% produce a document or report and 11.1% produce action artifacts—four times the action rate of any civic topic. Here people genuinely are using Claude to *do* something with government rather than to learn about it.

This is what resolves the puzzle that government filings and public records appear at magnitudes comparable to elections. Two different things are going on. Some of it is an artifact of the taxonomy: “Government filings” is a Minor topic (0.06%) and “Elections” a Detailed one, so their raw shares were never like-for-like. But the substantive point survives the correction: administrative engagement with government is real, is roughly the scale of electoral interest, and—unlike electoral interest—involves actual production.

Crucially, that production is not *representation*. Filing a form, claiming a benefit, or navigating an agency is the citizen acting as a client of the state, not as a principal directing an agent. In Hall’s terms it belongs to the service-delivery strand of the information layer—government working better for the people who deal with it—rather than to the representation layer, which concerns whether citizens’ interests are conveyed faithfully into political decisions. Nothing in these data looks like drafting comments to agencies, contacting legislators, or coordinating with neighbors. The representation layer remains close to empty even in the one topic where people are visibly doing things.

The governance layer is, unsurprisingly, invisible here, and would be even if it were thriving: deliberation over the rules governing AI systems is not something one would expect to observe

in the topic distribution of consumer chat.

The takeaway. Present-day political usage of Claude is concentrated in the information layer, and splits within it into a large civic-informational strand (learning about politics and elections) and a smaller administrative strand (transacting with government) that is the only place action-shaped output appears at all. The representation layer, where AI would act on citizens’ behalf, is not yet visible in consumer usage. If the information layer is where the technology is arriving first, that is where its near-term democratic effects—good and bad—will be concentrated.

Two cautions before this is read as a verdict. First, the AEI covers Claude’s consumer surfaces; agentic and API traffic, where representation-layer activity would more plausibly live, is either excluded or reported only globally. Second, an artifact label captures the most prominent output of a conversation, so a conversation that informed someone who later acted is coded as information. What we can say is narrower than “AI is not being used for political action”: it is that within consumer Claude usage, informational political activity currently outweighs action-shaped political activity by more than an order of magnitude. If the information layer is where the technology is currently arriving, that is where its near-term democratic effects—good and bad—will be concentrated.

3. Do Primaries Move Political Conversation?

3.1. Design

We compare states holding a statewide primary in May 2026 to states with later primaries, before and after. For state s in month t ,

$$y_{st} = \beta (\text{Primary}_s \times \text{May}_t) + \alpha_s + \delta_t + \varepsilon_{st}, \quad (1)$$

where y_{st} is the politics-topic share, α_s and δ_t are state and month fixed effects, and standard errors are clustered by state. The panel is balanced across April and May. States holding March primaries (TX, NC, IL) are already treated and are dropped throughout. The estimation sample is 30 states, 8 of them treated.

3.2. Main result

Holding a May primary raises the politics share by 0.10 percentage points (s.e. 0.04, $t = 2.61$) on a control-group base of 0.48%—an increase of about 22%. In levels this moves a state from roughly one political conversation in 250 to one in 200.

Table 5: Claude Conversations Shift Toward Politics When a State Holds Its Primary.

	Politics-Topic Share (%)		Log Share
	(1)	(2)	(3)
Primary $\times$ May	0.10 (0.04)	0.09 (0.04)	0.24 (0.09)
# States	30	20	30
# Treated States	8	8	8
# Observations	60	40	60
Control May Mean	0.48	0.45	0.48
Minimum Detectable Effect	0.11	0.12	0.24
State Fixed Effects	Yes	Yes	Yes
Month Fixed Effects	Yes	Yes	Yes
Excl. June-Primary Controls	No	Yes	No

Notes: Each column reports a difference-in-differences estimate comparing the politics-topic share of Claude conversations in April and May 2026 between states holding a statewide primary in May 2026 and states with later primaries. The outcome in columns 1 and 2 is the percent of a state’s conversations assigned to the “Politics and public record” topic in the Anthropic Economic Index; column 3 uses its natural log. Texas, North Carolina, and Illinois, which held March primaries, are excluded. Column 2 further drops control states with June primaries, whose mail voting begins in May. Standard errors clustered by state are in parentheses. The minimum detectable effect is 2.80 times the standard error.

3.3. The effect is specific to politics

If the estimate reflected a general election-season rise in information seeking, it should appear in adjacent topics too. It does not. Table 6 applies the identical specification to news aggregation, news writing, geopolitics, and a non-political comparison topic. All four are small and statistically indistinguishable from zero, and their standard errors are if anything tighter than the main estimate’s—so they are precise enough to have detected an effect the size of the politics result had one been present. Domestic politics moves; foreign affairs and general news do not.

Table 6: Specificity: The Same Design Applied to Adjacent Topics

Topic	Estimate	(SE)	t	# States	# Treated
Politics and public record	0.10	(0.04)	2.61	30	8
Geopolitics and strategy	-0.01	(0.02)	-0.39	16	4
Personal finance	-0.01	(0.03)	-0.44	25	6
News writing	-0.02	(0.03)	-0.55	16	4
News aggregation	-0.04	(0.03)	-1.17	26	6

3.4. A permutation test across all topics

The sharpest test available is to apply the identical design to every topic in the release and ask how unusual the politics estimate is. Among the 210 topic series with sufficient state coverage, politics ranks 3rd by t -statistic, and 3 topics have a t at least as large in absolute value. That yields a randomization-inference p -value of 0.024 without relying on the clustered standard errors at all (Figure 4).




Figure 4: Empirical null distribution. Distribution of t -statistics on Primary $\times$ May across all 210 topic series estimable with the same specification.

Two features of this distribution deserve comment. It is centered essentially at zero (mean -0.03), which is reassuring: the design is not mechanically manufacturing effects. But its dispersion exceeds the nominal (s.d. 1.17 against 1.00 under correctly sized standard errors), and 21 of 210 topics—10%—clear $|t| \geq 1.96$ against the 5% expected by chance. Conventional standard errors are therefore modestly too small here, which is exactly why the permutation p -value is the number to report. Some unrelated topics do land above politics; with 210 tests that is expected, and the permutation test prices it in rather than being embarrassed by it.

4. Limitations

Two periods, so no pre-trends. With only April and May, parallel trends is assumed rather than tested. A third month would allow a genuine pre-trend check and is the single most valuable addition a future AEI release could make.

Treated states are lost to the balance requirement. Of 11 May-primary states, 8

survive into the estimation sample; 3 (KY, NE, WV) are dropped because they fall below the privacy threshold in at least one month. These are small states, so this is most likely a volume artifact, but the treatment group is 3/11 smaller than the calendar implies.

Coarse measurement. Shares are rounded to two decimals, so the effect is ten rounding increments wide.

No comparable data elsewhere. Replicating this design with Google or OpenAI data is currently impossible. Google’s AI & Economy ATLAS (July 2026) reports country-level data for a single two-week window and publishes no dataset. OpenAI Signals publishes monthly state data under CC-BY, but its topic taxonomy is functional (programming, cooking, editing) rather than subject-matter, and contains no political category; its only state-by-topic file is an annual average. The AEI is, for now, the only public source combining US state, calendar month, and subject-matter topic.